

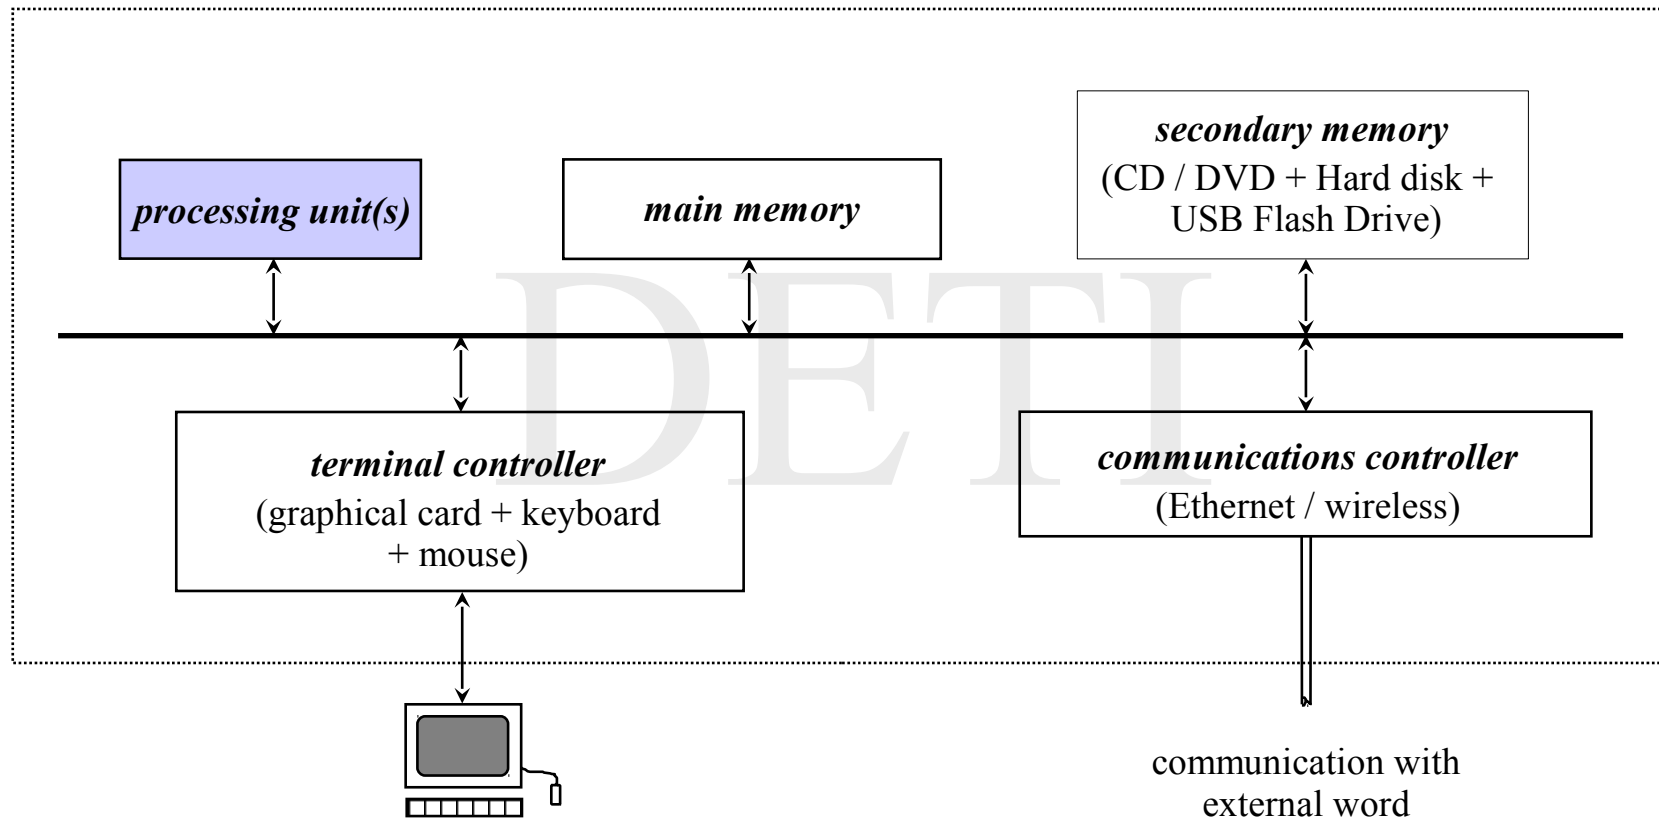


# *Operating Systems / Sistema de Operação*

*Processes and threads*

António Rui Borges / Artur Pereira

# *Typical computational system*



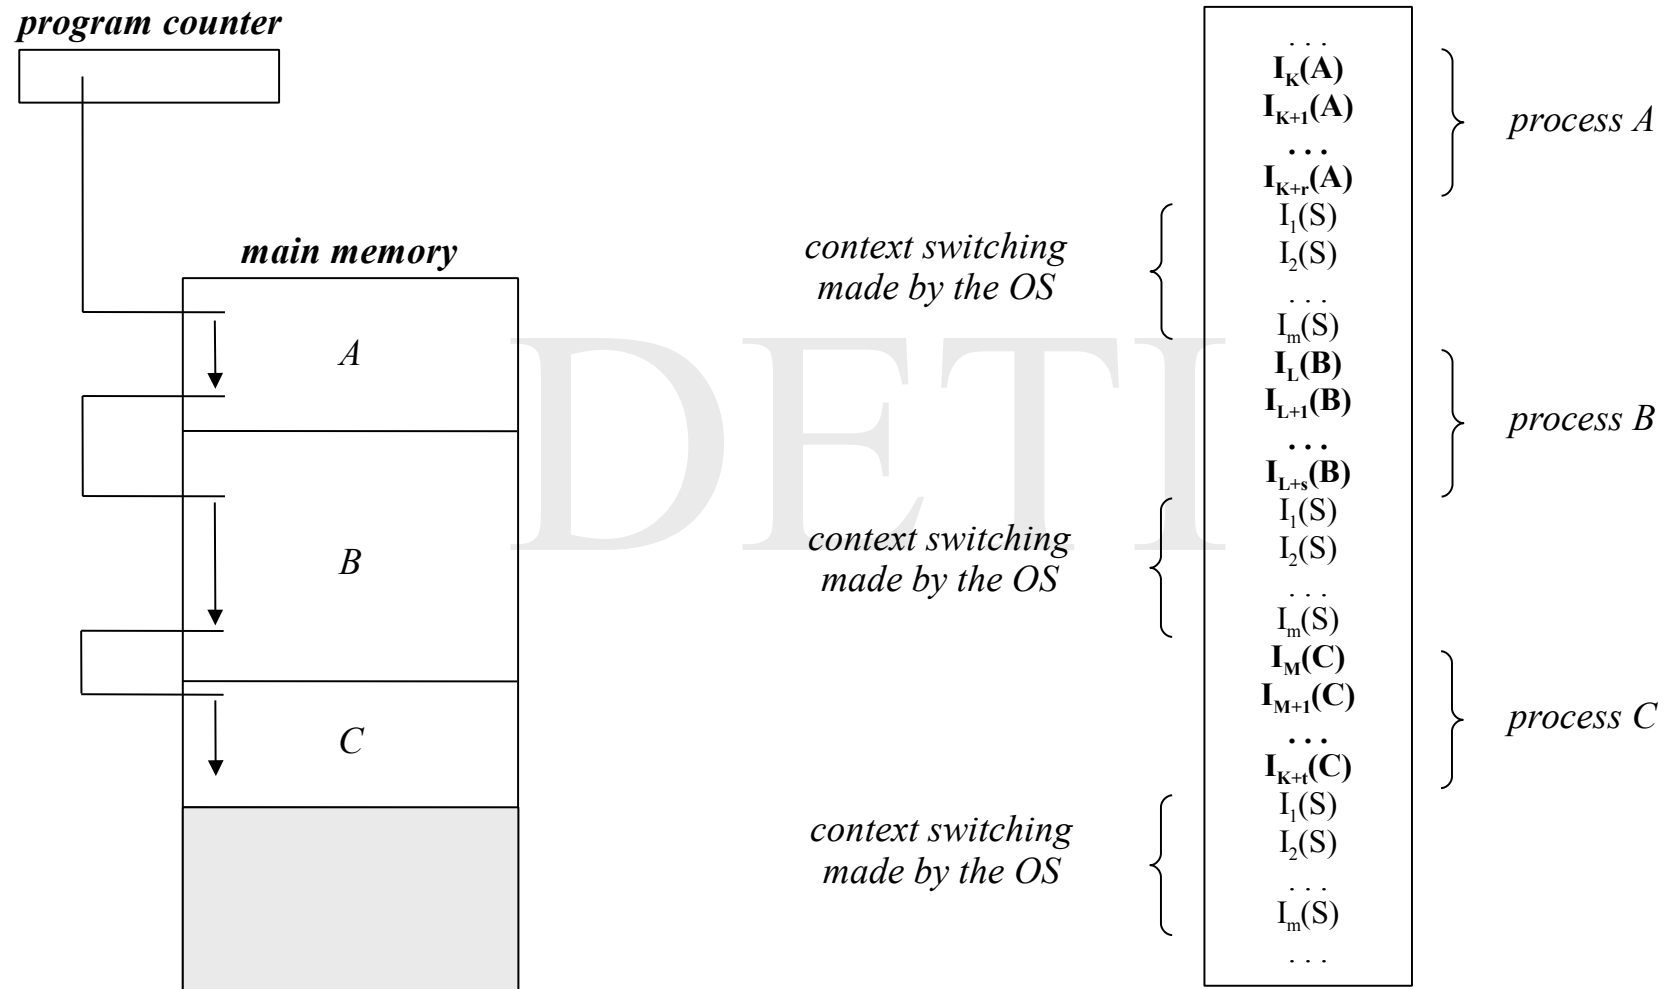
# *Contents*

- ♦ *Program vs. Process*
- ♦ *Process model*
- ♦ *State diagram of a process*
- ♦ *Process creation in Unix*
- ♦ *Address space of a Unix program*
- ♦ *Threads and multithreading*
- ♦ *Threads in Linux*
- ♦ *Process switching*
- ♦ *Bibliography*

## *Program vs. Process*


- ♦ **Program** – set of instructions describing how a task is performed by a computer
  - ♦ In order for the task to be actually performed, the corresponding program has to be executed
- ♦ **Process** – an entity that represents a computer program being executed
  - ♦ it represents an activity of some kind
  - ♦ it is characterized by:
    - ♦ code and data (actual values of the different variables) of the associated program (*addressing space*);
    - ♦ actual values of the processor internal registers
    - ♦ input and output data (data that are being transferred from input devices and to output devices)
    - ♦ state of execution
- ♦ Different processes can be running the same program
- ♦ In general, there are more processes than processors – *multiprogramming*

# Execution in a multiprogrammed environment




## *Process model*

- In **multiprogramming**, the activity of the processor, because it is switching back and forth from process to process, is hard to perceive
- Thus, it is better to assume the existence of a number of **virtual processors**, one per existing process
  - Turning *off* one virtual processor and *on* another, corresponds to a process switching
  - number of active virtual processors  $\leq$  number of real processors



## *Process model*

- The switching between processes, and thus the switching between virtual processors, can occur for different reasons, possible not controlled by the running program
- Thus, to be viable, this process model requires that
  - *the execution of any process is not affected by the instant in time or the location in the code where the switching takes place*
  - *no restrictions are imposed on the total or partial execution times of any process*

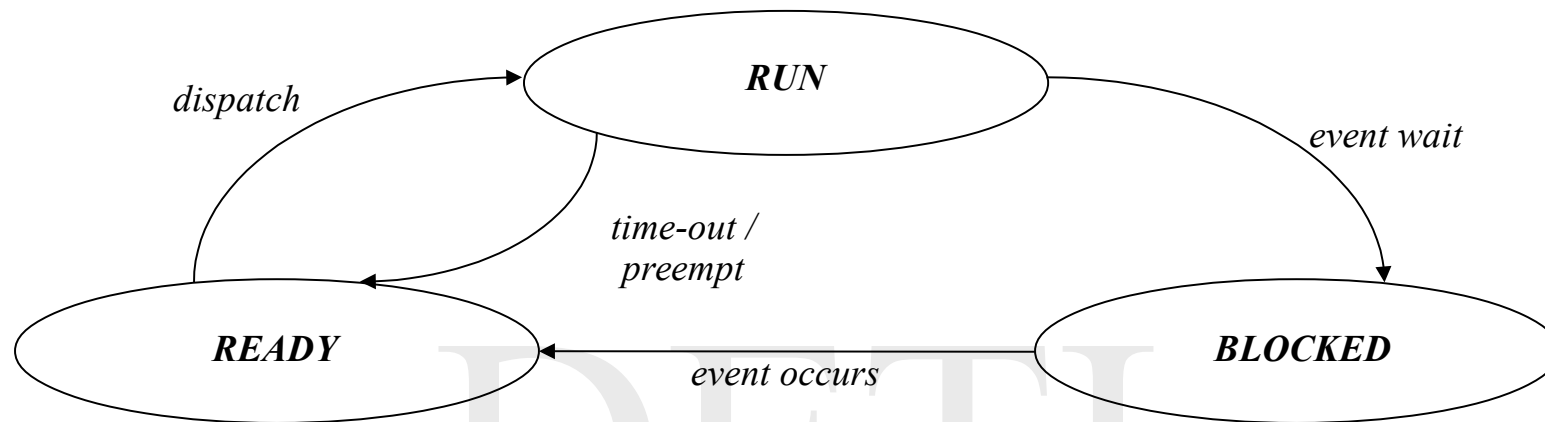


## *State diagram of a process*

- During its existence, a process can be in different situations, named *states*
- The most important are:
  - *run* – the process is in possession of a processor, and thus running
  - *blocked* – the process is waiting for the occurrence of an external event (access to a resource, end of an input/output operation, etc.)
  - *ready* – the process is ready to run, but waiting for the availability of a processor to start/resume its execution
- Transitions between states usually result from external intervention, but can in some cases be triggered by the process itself
- The part of the operating system that handles these transitions is called the (*processor*) *scheduler*, and is an integral part of its kernel
- Different policies exist to control the firing of these transitions



## State diagram of a process

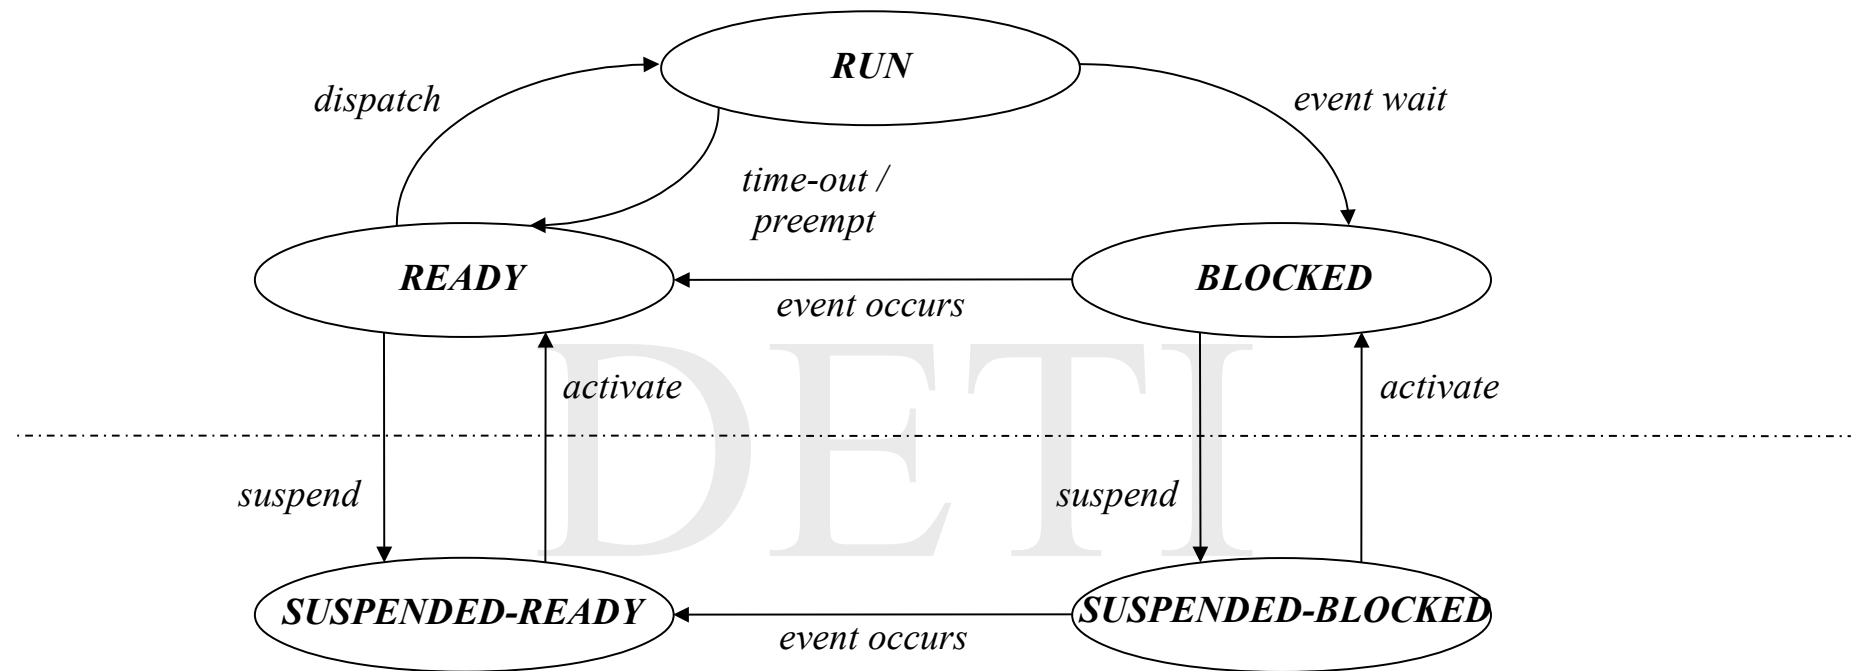


- *dispatch* – one of the processes ready to run is selected and is given the processor
- *event wait* – the running process is prevented to proceed, awaiting the occurrence of an external event
- *event occurs* – the external event occurred and the process must now wait for the processor
- *time-out* – the time slot assigned to the running process get to the end, so the process is removed from the processor
- *preempt* – a higher priority process get ready to run, so the running process is removed from the processor

## *State diagram of a process*

- The main memory is finite, which limits the number of coexisting processes.
- A way to overcome this limitation is to use an area in secondary memory to extend the main memory
  - This is called *swap area* (can be a disk partition or a file)
  - A non running process, or part of it, can be *swapped out*, in order to free main memory for other processes
  - That process will be later on *swapped in*, after main memory becomes available
- Two new states should be added to the process state diagram to incorporate these situations:
  - *suspended-ready* – the process is *ready* but swapped out
  - *suspended-blocked* – the process is *blocked* and swapped out

## State diagram of a process

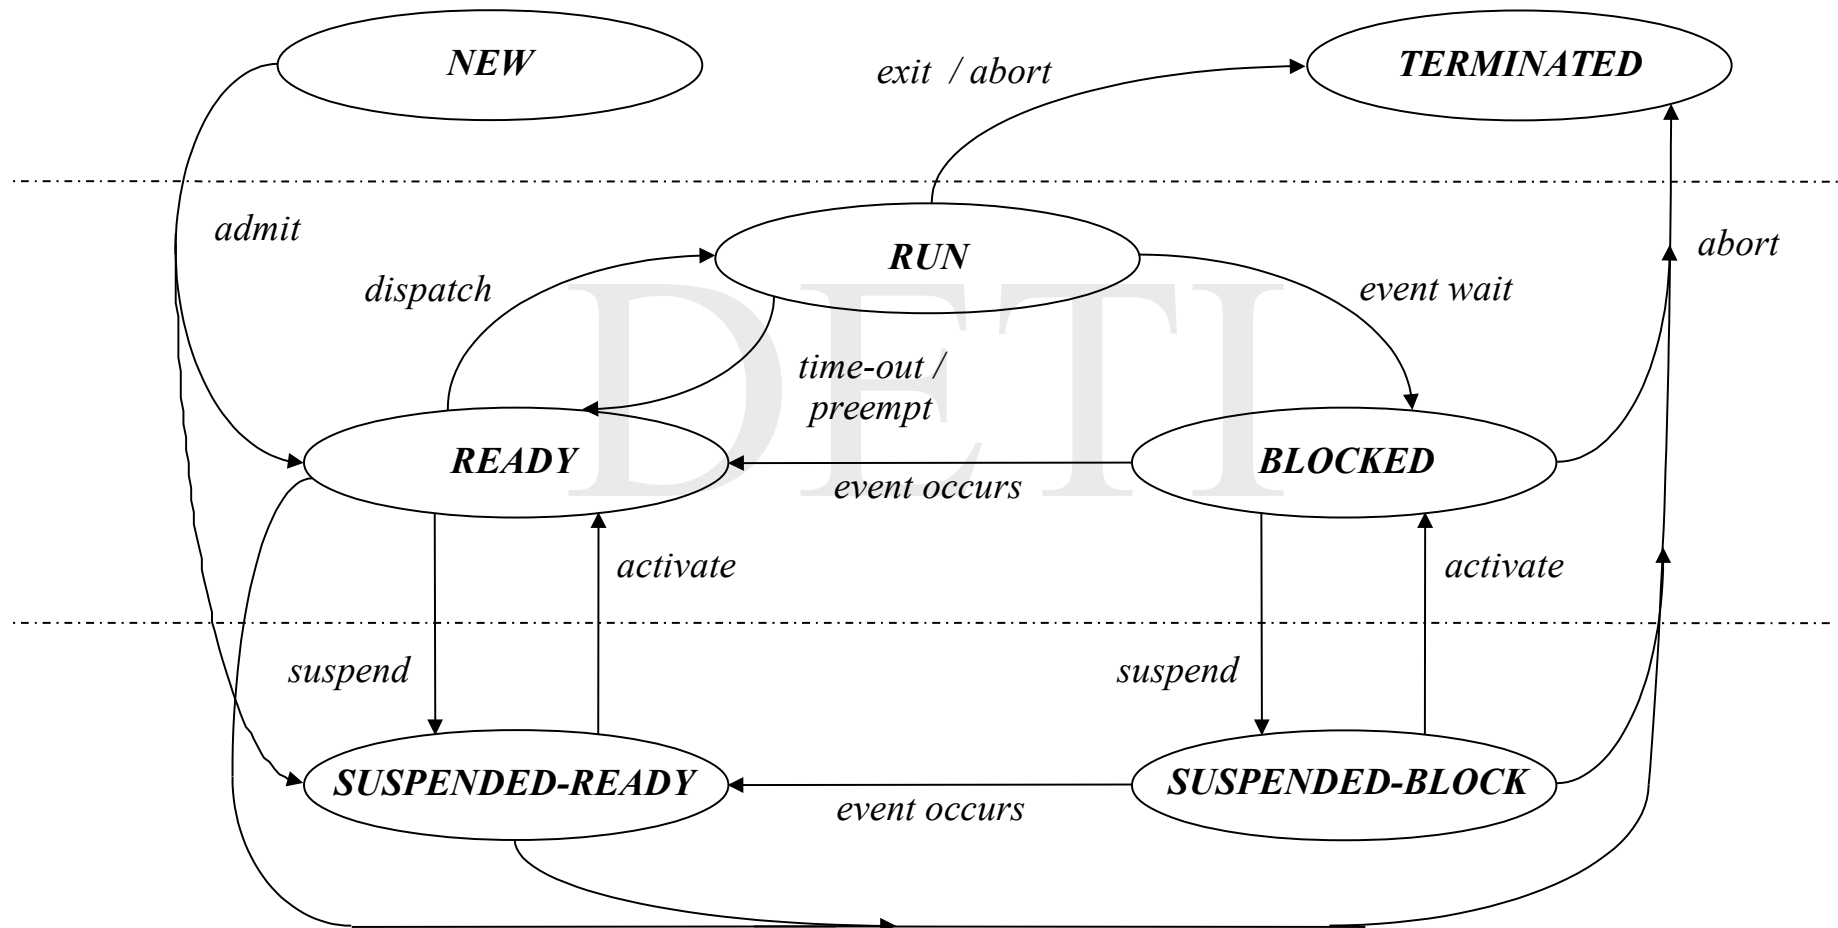


- ♦ Two new transitions appear:
  - ♦ *suspend* – the process is *swapped out*
  - ♦ *activate* – the process is *swapped in*


## *State diagram of a process*

- ♦ The previous state diagram assumes processes are timeless
  - ♦ Apart from some system processes this is not true
  - ♦ Processes are created, exist for some time, and eventually finish
- ♦ Two new states are required to represent creation and termination
  - ♦ *new* – the process has been created but not yet admitted to the pool of executable processes (the process data structure is been initialized)
  - ♦ *terminated* – the process has been released from the pool of executable processes (some actions are still required before the process is discarded)
- ♦ three new transitions exist
  - ♦ *admit* – the process is admitted (by the OS) to the pool of executable processes
  - ♦ *exit* – the running process indicates the OS it has completed
  - ♦ *abort* – the process is forced to terminate (because of a fatal error or because an authorized process aborts its execution)

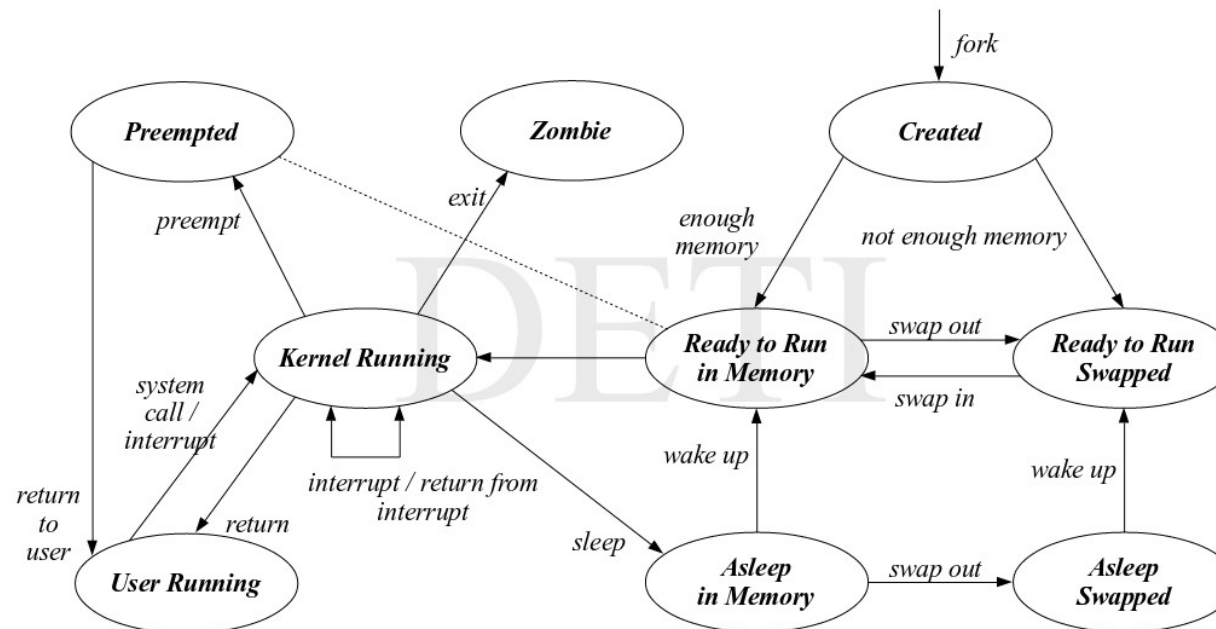
## State diagram of a process



## State diagram of a Unix process




## State diagram of a Unix process



- There are two *run* states, *kernel running* and *user running*, associated to the processor running mode, supervisor and user, respectively
- The *ready* state is also splitted in two states, *ready to run in memory* and *preempted*, but they are equivalent, as indicated by the dashed line

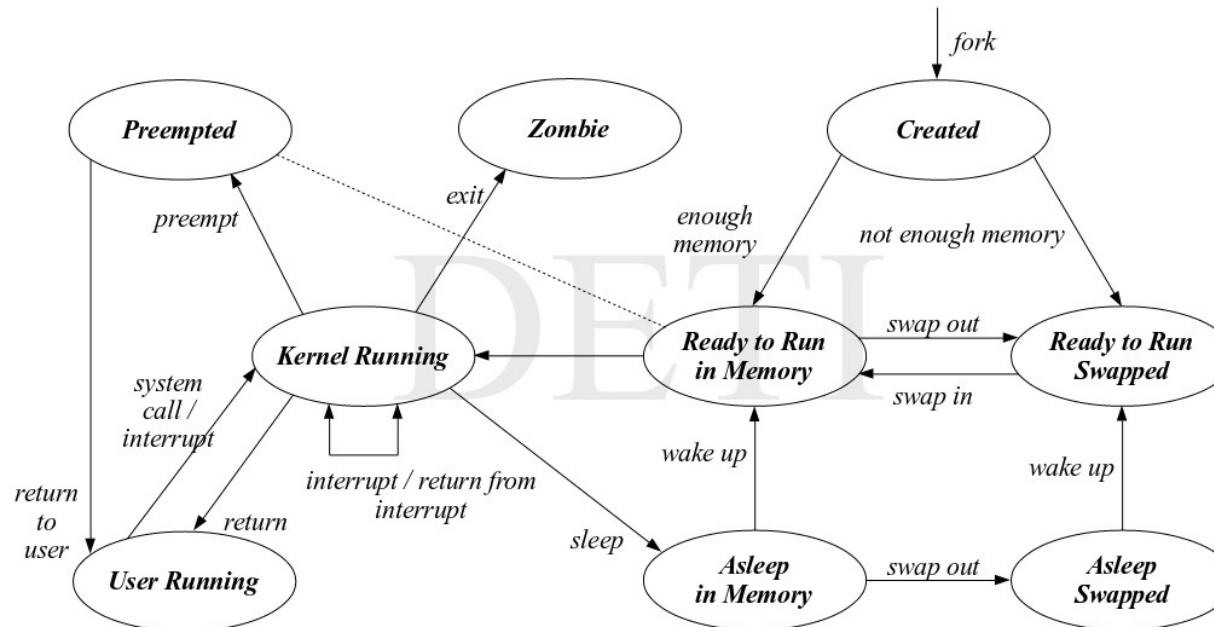
## State diagram of a Unix process



- When a user process leaves supervisor mode, it can be preempted (because a higher priority process is ready to run)
- In practice, processes in *ready to run in memory* and *preempted* shared the same queue, thus are treated as equal
- The *time-out* transition is covered by the *preempt* one

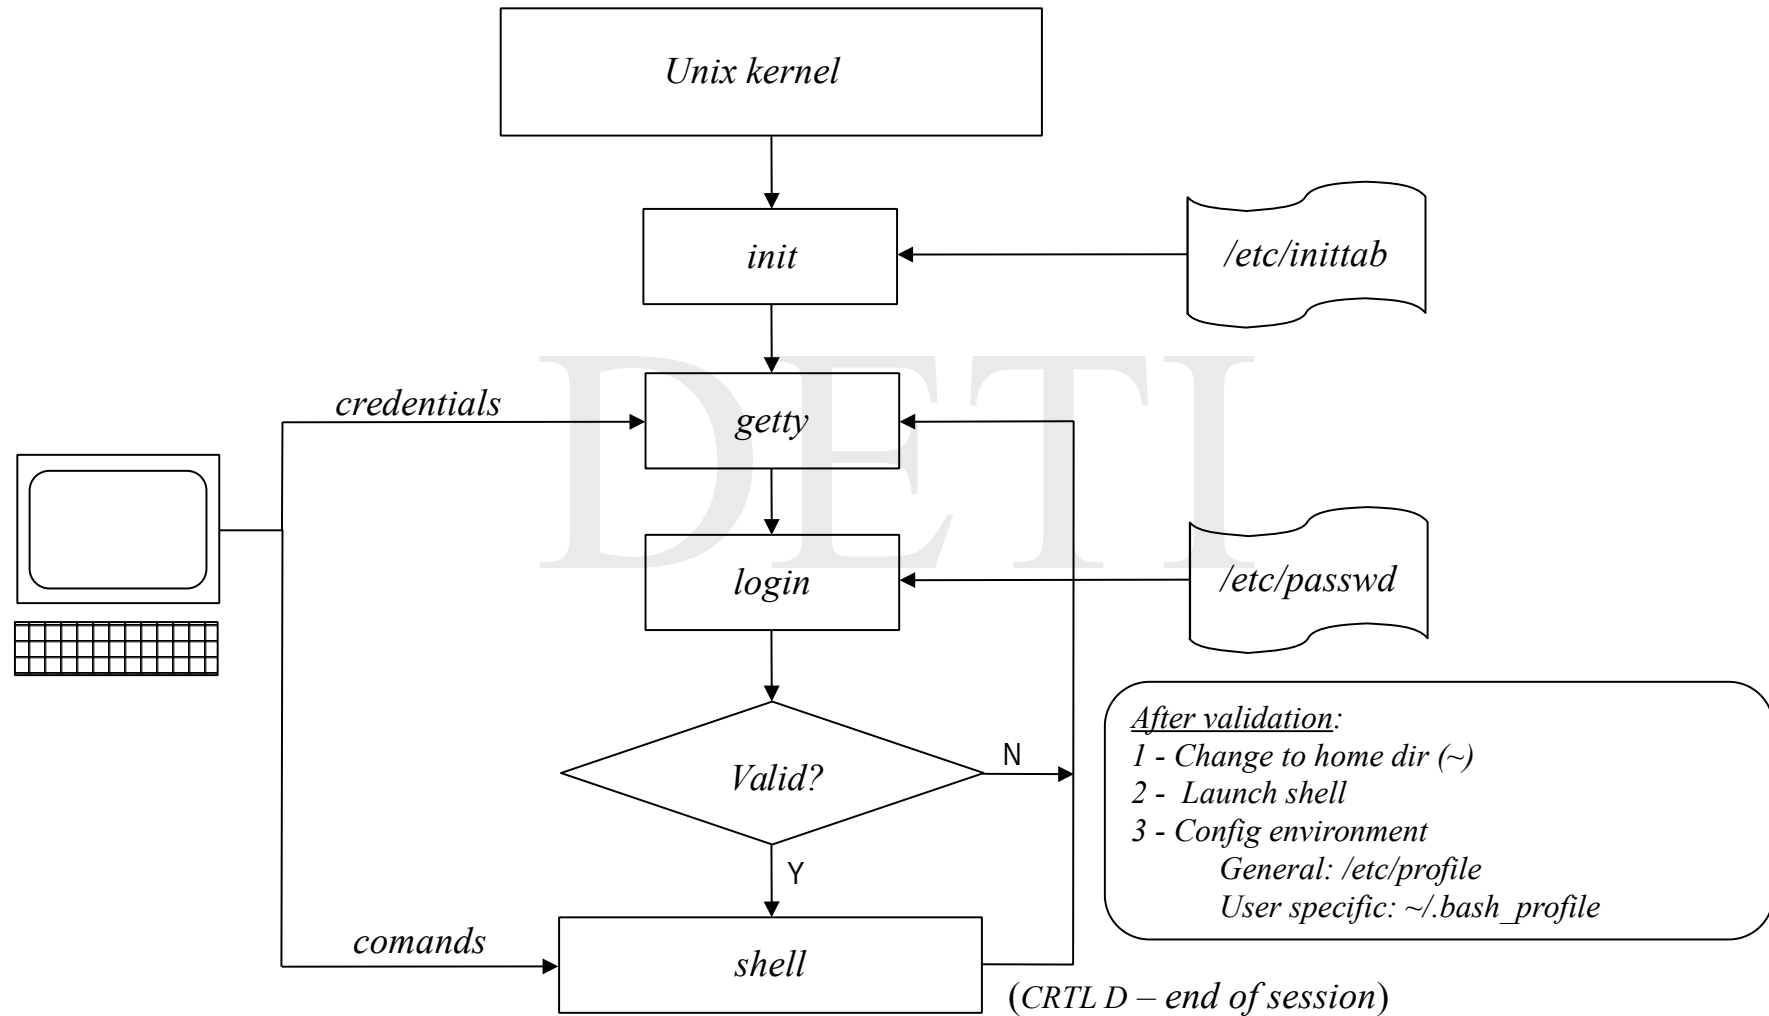


## State diagram of a Unix process

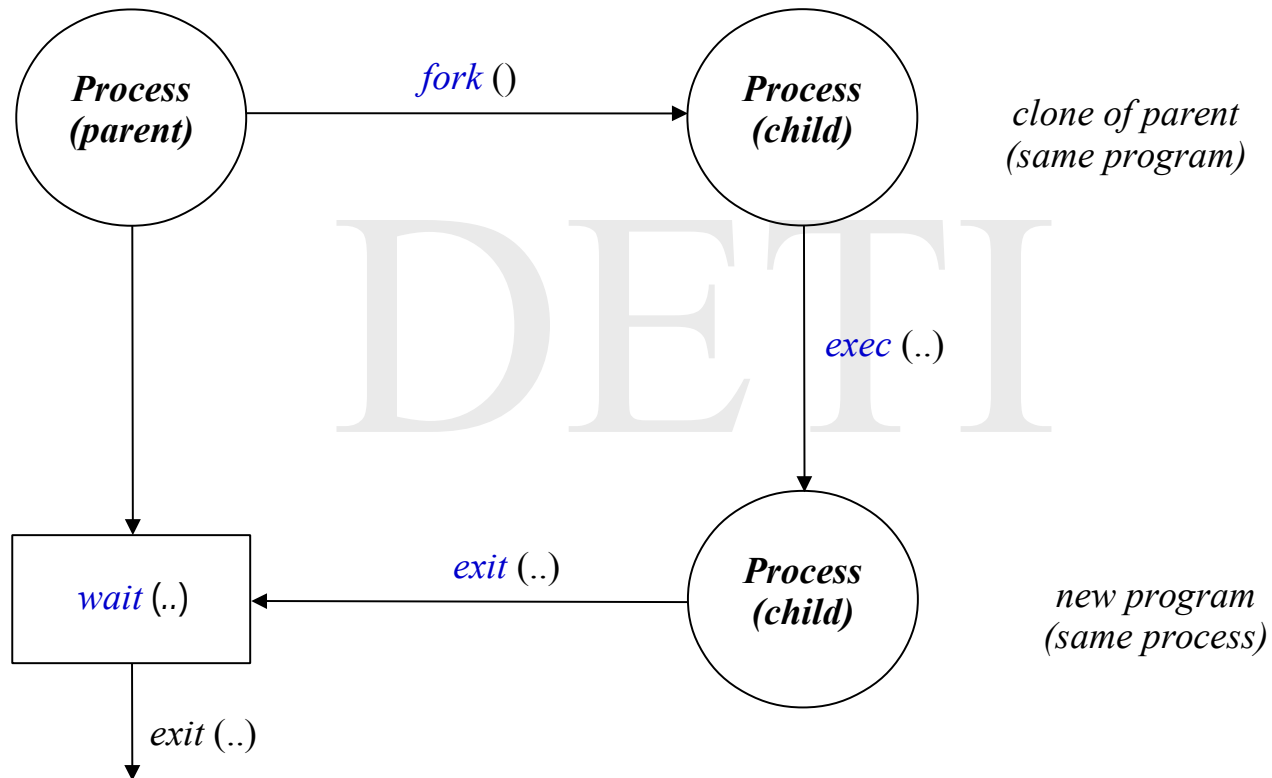


- traditionally, execution in supervisor mode could not be interrupted (thus UNIX does not allow real time processing)
- in current versions, namely from SVR4, the problem was solved by dividing the code into a succession of atomic regions between which the internal data structures are in a safe state and therefore allowing execution to be interrupted
- this corresponds to a transition between the *preempted* and *kernel running* states, that could be called *return to kernel*.

# Unix – traditional login



## *Process creation in Unix*



## *Process creation in Unix*

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>

int main(void)
{
    printf("Before the fork:\n");
    printf("  PID = %d, PPID = %d.\n",
        getpid(), getppid());

    fork();

    printf("After the fork:\n");
    printf("  PID = %d, PPID = %d. Who am I?\n",
        getpid(), getppid());

    return EXIT_SUCCESS;
}
```

- The **fork** clones the executing process, creating a replica of it
- The address spaces of the two processes are equal
  - actually, just after the fork, they are the same
  - typically, a copy on write approach is followed
- The states of execution are the same
  - including the program counter
- Some process variables are different (PID, PPID, ...)

## *Process creation in Unix*

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>

int main(void)
{
    printf("Before the fork:\n");
    printf("  PID = %d, PPID = %d.\n",
           getpid(), getppid());

    int ret = fork();

    printf("After the fork:\n");
    printf("  PID = %d, PPID = %d, ret = %d\n",
           getpid(), getppid(), ret);

    return EXIT_SUCCESS;
}
```

- The value returned by the fork is different in parent and child
  - in parent is the PID of the child
  - in child is always 0
- This return value can be used as a boolean variable, so we can distinguish the code running on child and parent

## *Process creation in Unix*

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>

int main(void)
{
    printf("Before the fork:\n");
    printf("  PID = %d, PPID = %d.\n",
           getpid(), getppid());

    int ret = fork();

    if (ret == 0)
    {
        printf("I'm the child:\n");
        printf("  PID = %d, PPID = %d\n",
               getpid(), getppid());
    }
    else
    {
        printf("I'm the parent:\n");
        printf("  PID = %d, PPID = %d\n",
               getpid(), getppid());
    }

    return EXIT_SUCCESS;
}
```

- By itself, the fork is of little interest
- In general, we want to run a different program in the child
  - `exec` system call
  - there are different versions of exec
- Sometimes, we want the parent to wait for the conclusion of the program running in the child
  - `wait` system call

# *Process creation in Unix*

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/wait.h>

int main(int argc, char *argv[])
{
    /* check arguments */
    if (argc != 2)
    {
        fprintf(stderr, "spawn <path to file>\n");
        exit(EXIT_FAILURE);
    }
    char *aplic = argv[1];

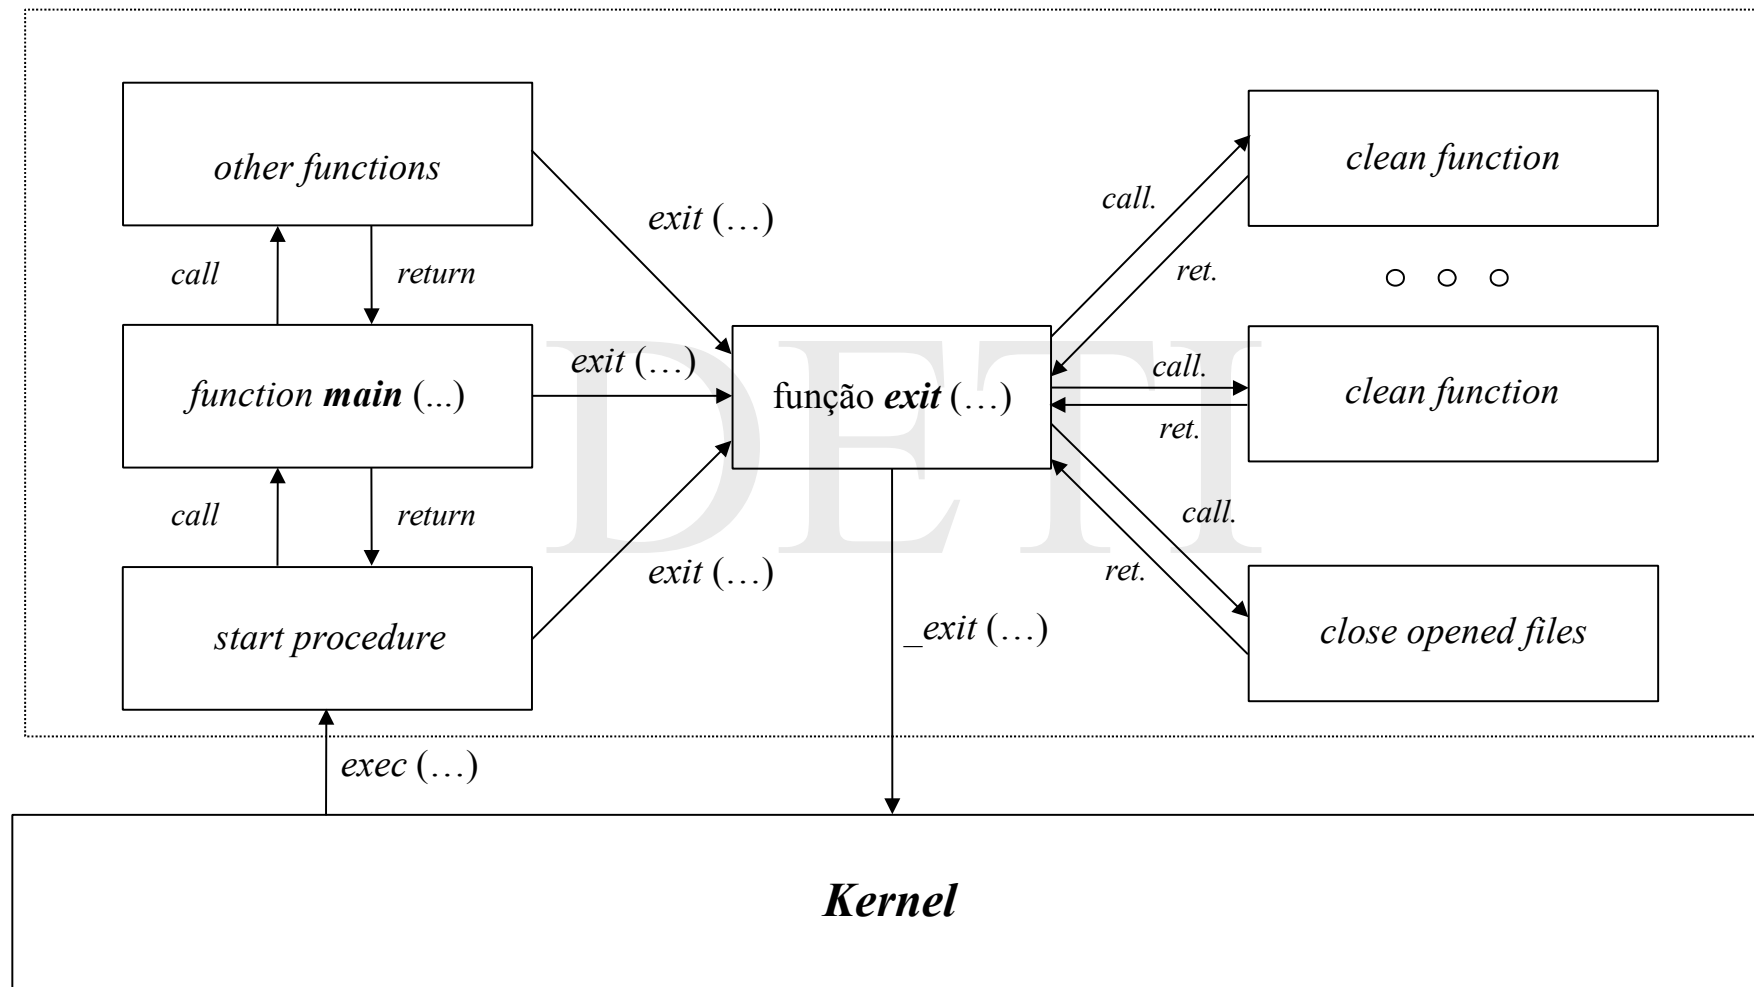
    printf("=====\n");

    /* clone phase */
    int pid;
    if ((pid = fork()) < 0)
    {
        perror("Fail cloning process");
        exit(EXIT_FAILURE);
    }
}
```

```
/* exec and wait phases */
if (pid != 0) // only runs in parent process
{
    int status;
    while (wait(&status) == -1);
    printf("=====\n");
    printf("Process %d (child of %d)"
           " ends with status %d\n",
           pid, getpid(), WEXITSTATUS(status));
}
else // this only runs in the child process
{
    execl(aplic, aplic, NULL);
    perror("Fail launching program");
    exit(EXIT_FAILURE);
}

exit(EXIT_SUCCESS); // or return EXIT_SUCCESS
}
```

## Executing a C/C++ program





## *Executing a C/C++ program*

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <assert.h>

/* cleaning functions */
static void atexit_1(void)
{
    printf("atexit 1: %d\n", ++a);
}

static void atexit_2(void)
{
    printf("atexit 2: %d\n", ++a);
}

/* programa principal */
int main(void)
{
    /* registering at exit functions */
    assert(atexit(atexit_1) == 0);
    assert(atexit(atexit_2) == 0);

    /* normal work */
    printf("hello world 1!\n");

    for (int i = 0; i < 5; i++) sleep(1);

    return EXIT_SUCCESS;
}
```

- The `atexit` function allows to register a function to be called at the program's normal termination
- They are called in reverse order relative to their register
- What happens if the termination is forced?

# *Command line arguments and environment variables*

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[], char *env[])
{
    /* printing command line arguments */
    printf("Command line arguments:\n");
    for (int i = 0; argv[i] != NULL; i++)
    {
        printf("  %s\n", argv[i]);
    }

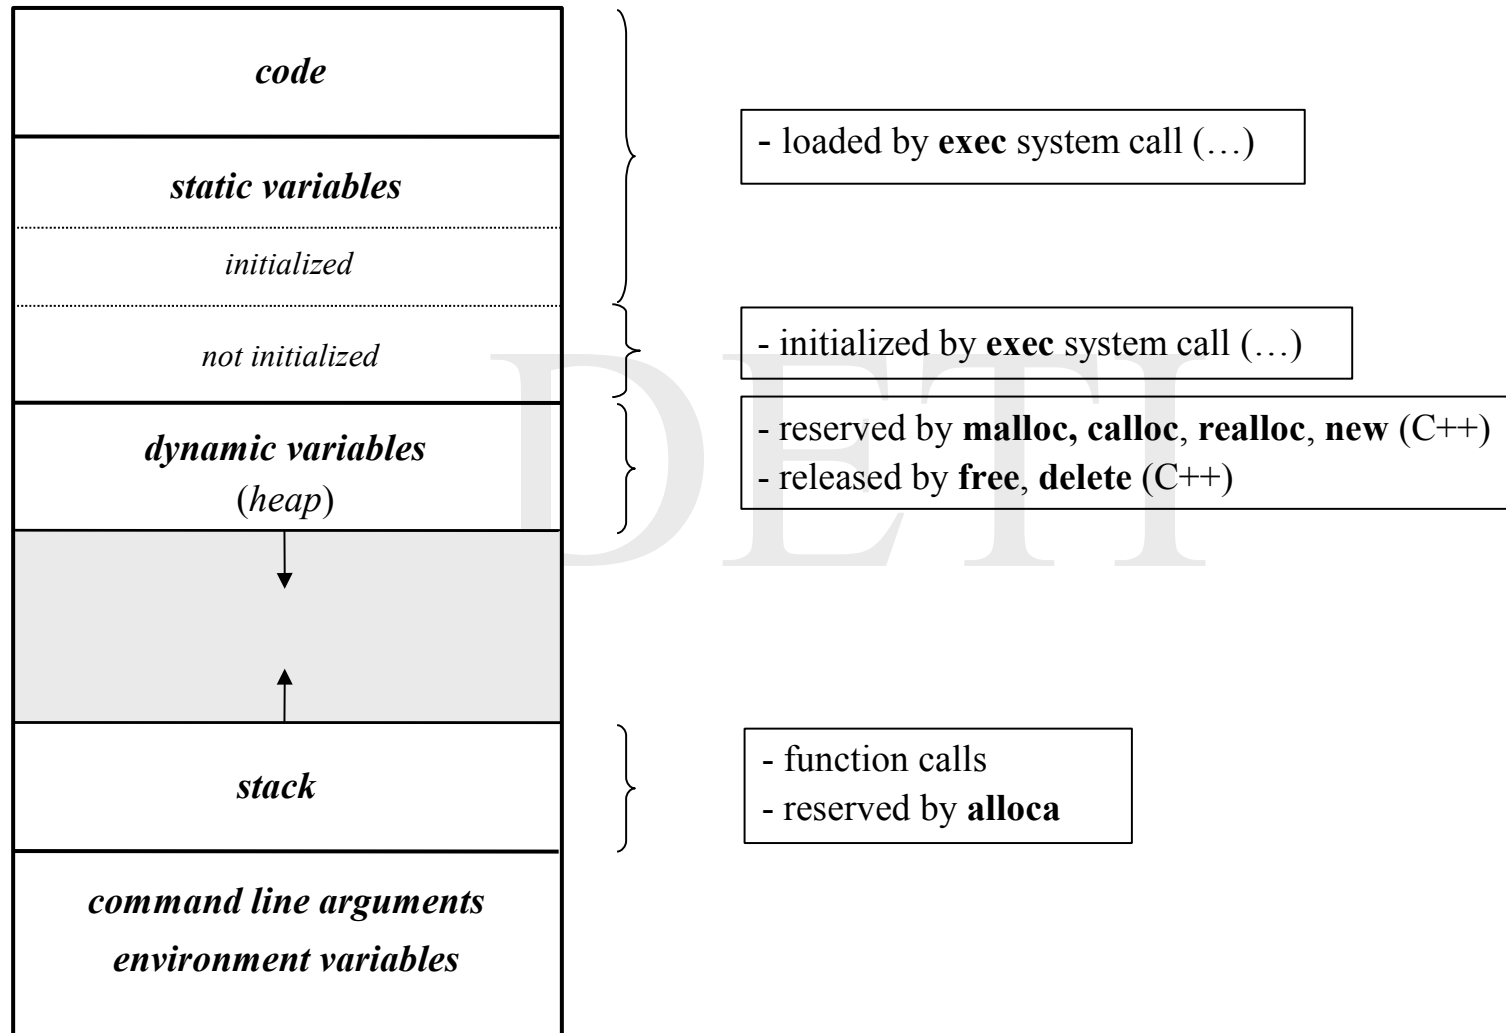
    /* printing all environment variables */
    printf("\nEnvironment variables:\n");
    for (int i = 0; env[i] != NULL; i++)
    {
        printf("  %s\n", env[i]);
    }

    /* printing a specific environment variable */
    printf("\nEnvironment variable:\n");
    printf("  env[\"HOME\"] = \"%s\"\n", getenv("HOME"));
    printf("  env[\"zzz\"] = \"%s\"\n", getenv("zzz"));

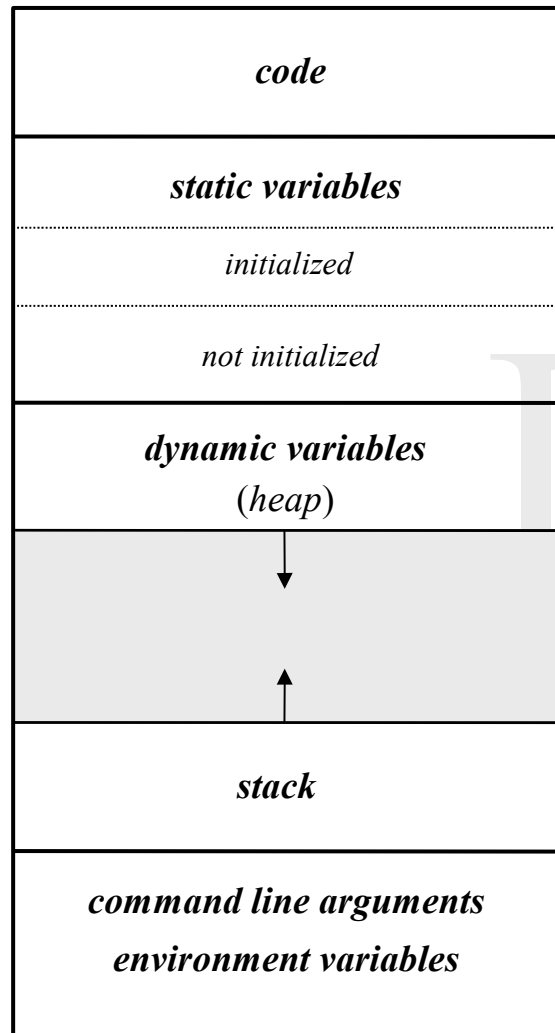
    return EXIT_SUCCESS;
}
```

- `argv` is an array of strings
  - `argv[0]` is the program reference
- `env` is an array of strings, each one representing a variable, in the form **name-value** pair
- `getenv` returns the value of a variable name

# Address space of a Unix process



# Address space of a Unix process



```


int n1 = 1;
static int n2 = 2;
int n3;
static int n4;
int n5;
static int n6 = 6;

int main(int argc, char *argv[], char *env[])
{
    extern char** environ;
    static int n7 = 3;
    static int n8;
    int *n9 = (int*)malloc(sizeof(int));
    int *n10 = (int*)malloc(sizeof(int));
    int *n11 = (int*)alloca(sizeof(int));
    int n12;
    int n13 = 13;
    int n14;
    printf("\ngetenv(n0) = %p\n", getenv("n0"));
    printf("\nargv = %p\nenviron = %p\nenv = %p\nmain = %p\n\n",
           argv, environ, env, main);
    printf("\n&argc = %p\n&argv = %p\n&env = %p\n",
           &argc, &argv, &env);
    printf("&n1 = %p\n&n2 = %p\n&n3 = %p\n&n4 = %p\n&n5 = %p\n"
           "&n6 = %p\n&n7 = %p\n&n8 = %p\n&n9 = %p\n&n10 = %p\n"
           "&n11 = %p\n&n12 = %p\n&n13 = %p\n&n14 = %p\n",
           &n1, &n2, &n3, &n4, &n5, &n6, &n7, &n8,
           &n9, &n10, &n11, &n12, &n13, &n14);

    return EXIT_SUCCESS;
}

```

# Address space of a Unix process



```

int n1 = 1;
...
static int n6 = 6;

int main(int argc, char *argv[], char *env[])
{
    if (fork() != 0)
    {
        fprintf(stderr, "n1 = %d\n", n1);
        wait(NULL);
        fprintf(stderr, "=====\n");
        fprintf(stderr, "n1 = %d\n", n1);
        fprintf(stderr, "=====\n");
    }
    else
    {
        n1 = 1111;
        fprintf(stderr, "n1 = %d\n", n1);
    }

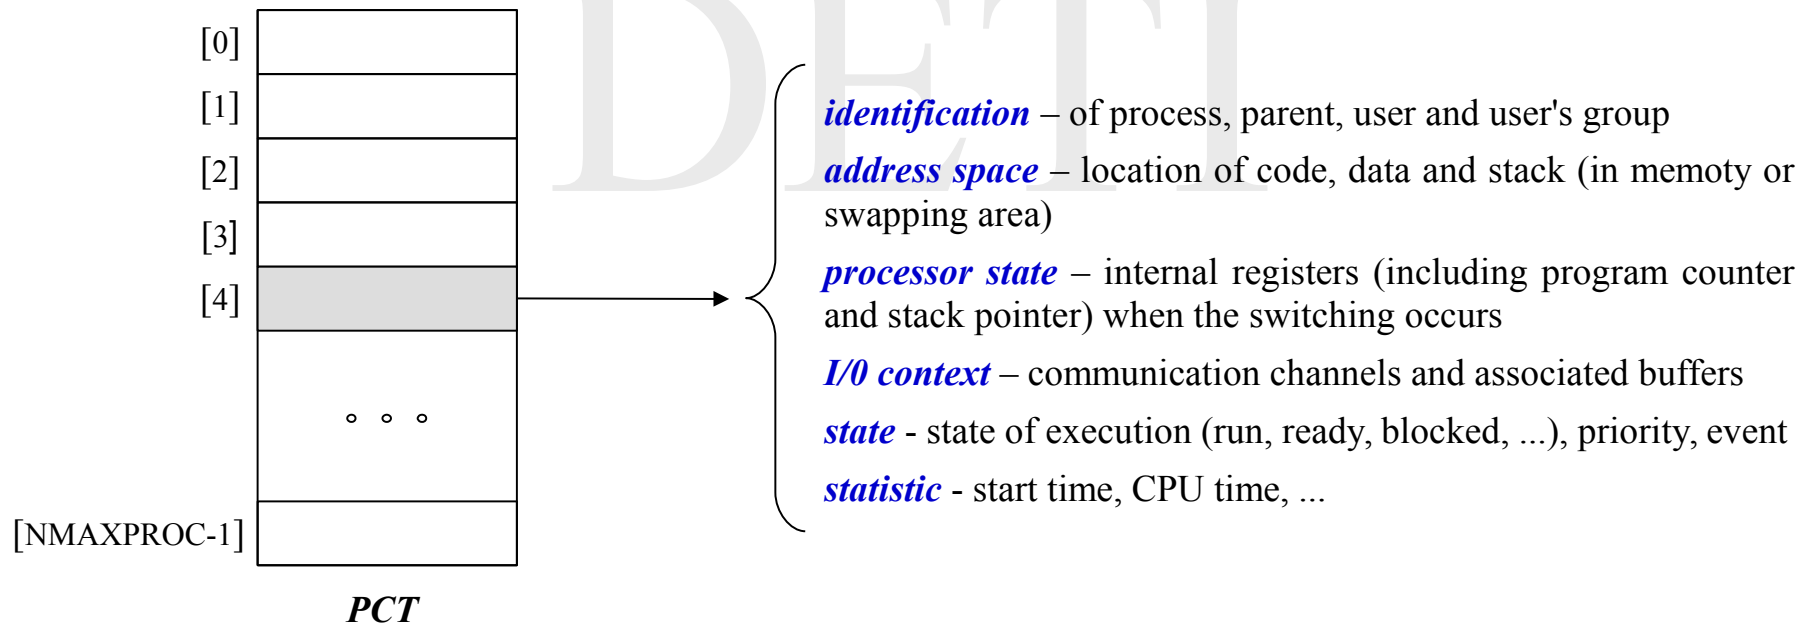
    static int n7 = 3;
    static int n8;
    ...
    printf("&n1 = %p\n&n2 = %p\n&n3 = %p\n&n4 = %p\n&n5 = %p\n"
        "&n6 = %p\n&n7 = %p\n&n8 = %p\n&n9 = %p\n&n10 = %p\n"
        "&n11 = %p\n&n12 = %p\n&n13 = %p\n&n14 = %p\n",
        &n1, &n2, &n3, &n4, &n5, &n6, &n7, &n8,
        &n9, &n10, &n11, &n12, &n13, &n14);

    return EXIT_SUCCESS;
}

```

## *Process control table*

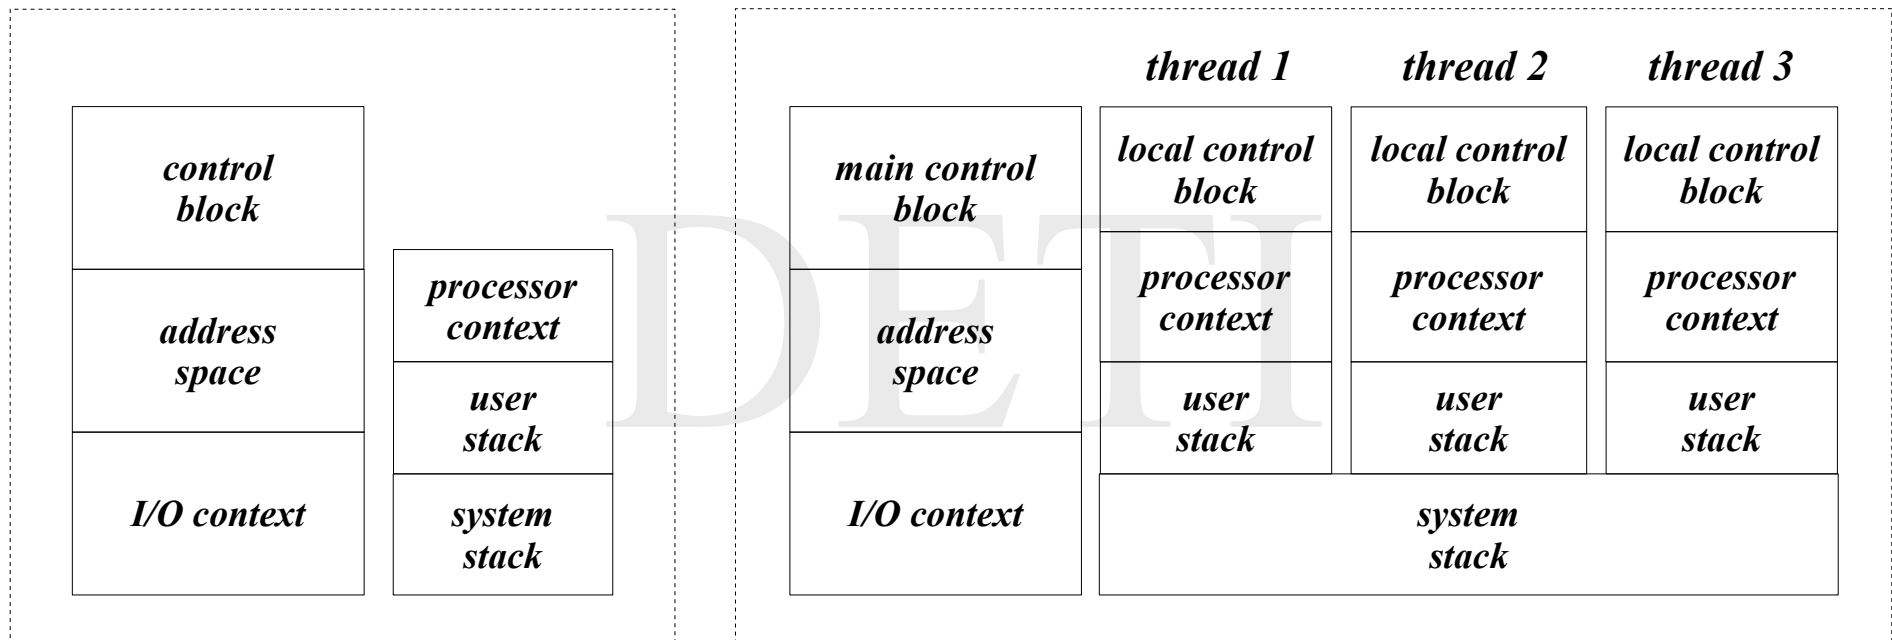
- To implement the process model, the operating systems needs a data structure to be used to store the information about each process – **process control block**
- The **process control table (PCT)**, an array of process control blocks, stores information about all processes



# Threads

- ♦ In traditional operating system, a process includes:
  - ♦ an address space and a set of communication channels with I/O devices
  - ♦ a single thread of control, which incorporates the processor registers (including the program counter) and a stack
- ♦ However, these components can be managed separately
- ♦ In this model, *thread* appears as an execution component within a process
- ♦ Several independent threads can coexist in the same process, thus sharing the same address space and the same I/O context
  - ♦ This is referred to as *multithreading*
- ♦ In practice, threads can be seen as *light weight processes*

# Multithreading

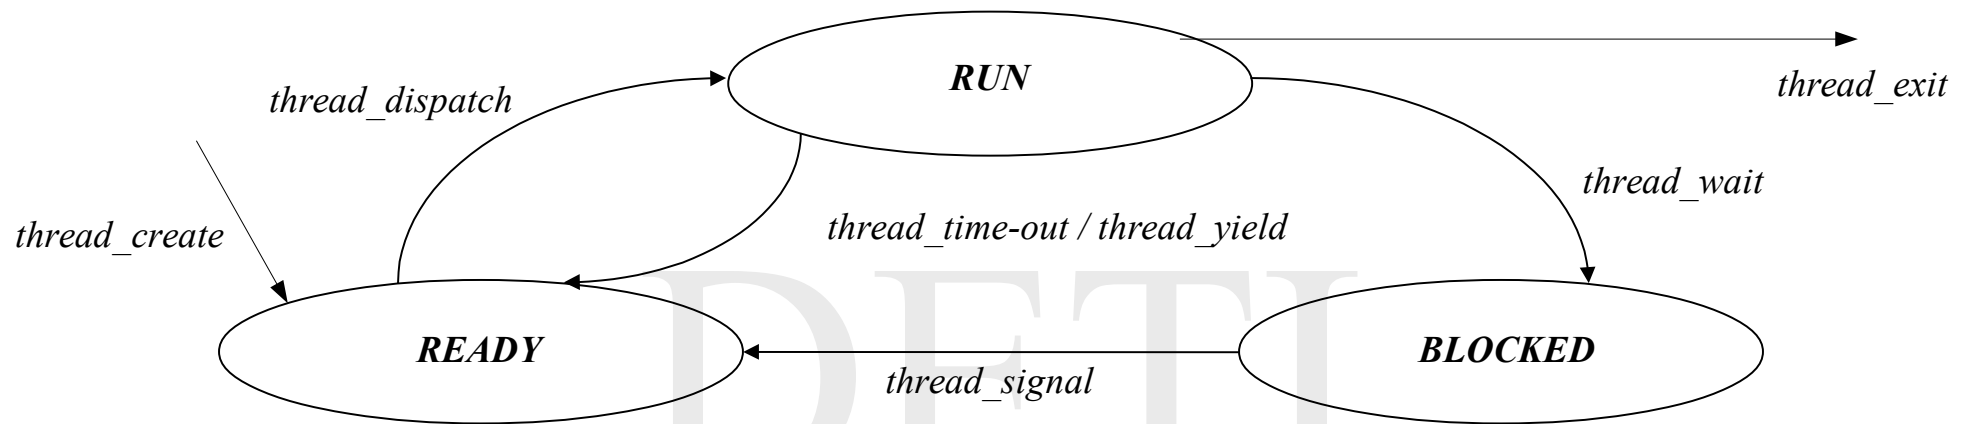


*Single threading*

*Multithreading*



## *State diagram of a thread*

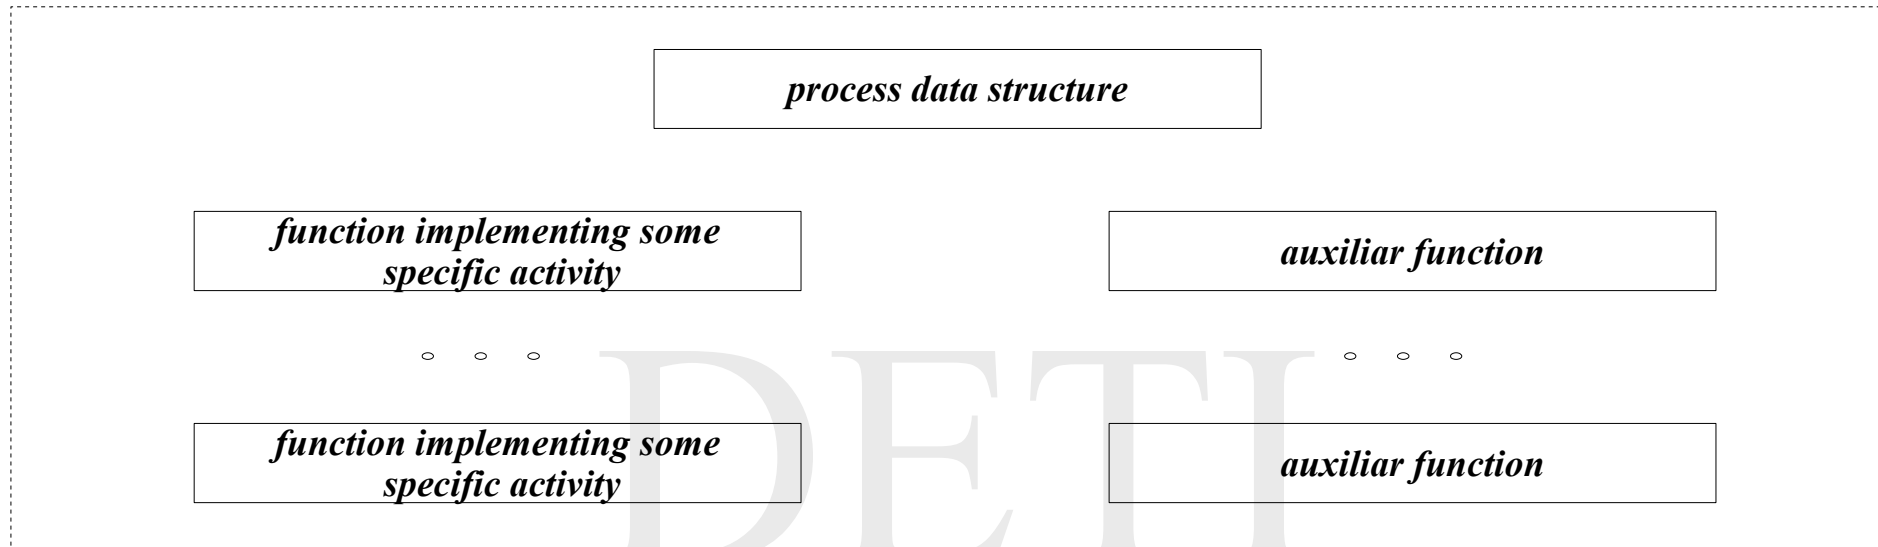


- only states concerning the management of the processor are considered
- states **SUSPENDED-READY** and **SUSPENDED-BLOCKED** are not present, since they are related to the address space and thus are related to the process, not to the threads
- states **NEW** and **TERMINATED** are not present, since the management of the multiprogramming environment is basically related to restrict the number of threads that can exist within a process

## *Advantages of multithreading*

- ♦ *easier implementation of applications* – in many applications, decomposing the solution into a number of parallel activities makes the programming model simpler
  - ♦ since the address space and the I/O context is shared among all threads, multithreading favors this decomposition.
- ♦ *better management of computer resources* – creating, destroying and switching threads is easier than doing the same with processes
- ♦ *better performance* – when an application involves substantial I/O, multithreading allows activities to overlap, thus speeding up its execution
- ♦ *multiprocessing* – real parallelism is possible if multiples CPUs exist.

## *Structure of a multithreaded program*

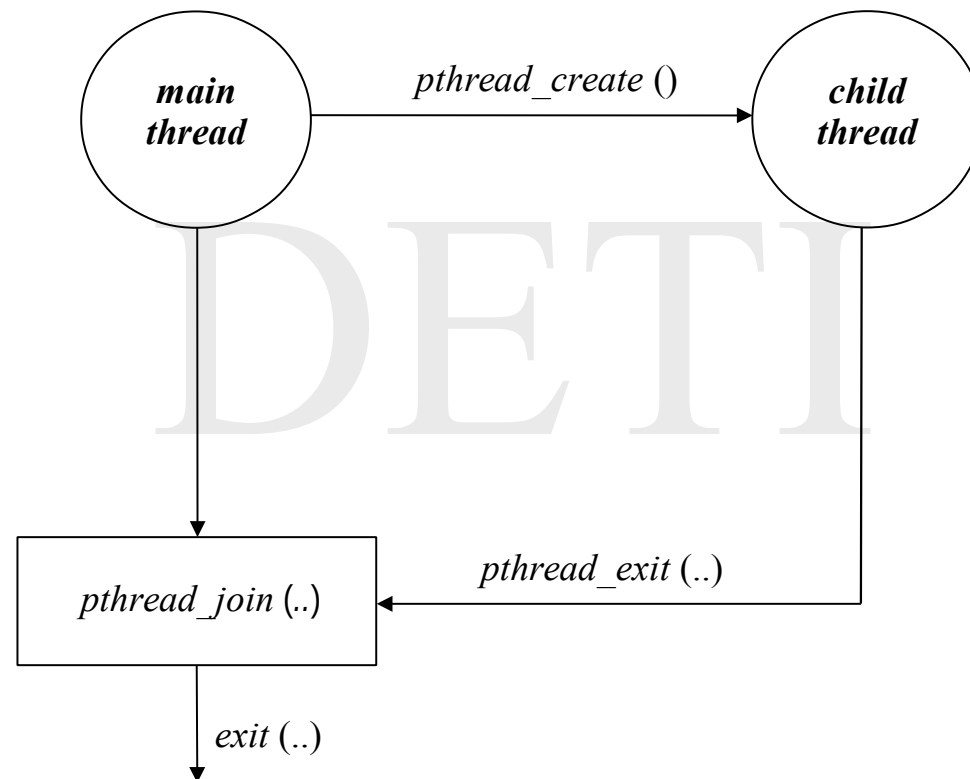


- each thread is typically associated to the execution of a function that implements some specific activity
- communication between threads can be done through the process data structure, which is global from the threads point of view
- the main program, also represented by a function that implements a specific activity, is the first thread to be created and, in general, the last to be destroyed

## *Implementations of multithreading*

- ♦ *user level threads* – threads are implemented by a library, at user level, which provides creation and management of threads without kernel intervention
  - ♦ versatile and portable
  - ♦ when a thread calls a blocking system call, the whole process blocks
    - ♦ kernel only sees the process
- ♦ *kernel level threads* – threads are implemented directly at kernel level
  - ♦ less versatile and portable
  - ♦ when a thread calls a blocking system call, another thread can be scheduled to execution

# *Library pthread*



## *Library pthread*

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

/* return status */
int status;

/* child thread */
void *threadChild (void *par)
{
    printf ("I'm the child thread!\n");

    status = EXIT_SUCCESS;
    pthread_exit (&status);
}
```

```
/* main thread */
int main (int argc, char *argv[])
{
    /* launching the child thread */
    pthread_t thr;
    if (pthread_create (&thr, NULL,
                       threadChild, NULL) != 0)
    {
        perror ("Fail launching thread");
        return EXIT_FAILURE;
    }

    /* waits for child termination */
    if (pthread_join (thr, NULL) != 0)
    {
        perror ("Fail joining child");
        return EXIT_FAILURE;
    }

    printf ("Child ends; status %d.\n", status);

    return EXIT_SUCCESS;
}
```

## *Threads in Linux*

- ♦ In Linux there are two system calls to create a child process:
  - ♦ *fork* – creates a new process that is a full copy of the current one
    - ♦ the address space and I/O context are duplicated
  - ♦ *clone* – creates a new process that can share elements with its parent
    - ♦ address space, table of file descriptors, and table of signal handlers, for example, are shareable.
    - ♦ the child starts execution in a given function
- ♦ Thus, from the kernel point of view, processes and threads are treated similarly
- ♦ Threads of the same process forms a thread group and have the same thread group identifier (TGID)
  - ♦ this is the value returned by system call `getpid()`
- ♦ Within a group, threads can be distinguished by their unique thread identifier (TID)
  - ♦ this value is returned by system call `gettid()`

# Threads in Linux

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <unistd.h>
#include <sys/syscall.h>
#include <sys/types.h>

pid_t gettid()
{
    return syscall(SYS_gettid);
}

/* child thread */
int status;
void *threadChild (void *par)
{
    printf ("Child: PPID: %d, PID: %d, TID: %d\n",
            getppid(), getpid(), gettid());

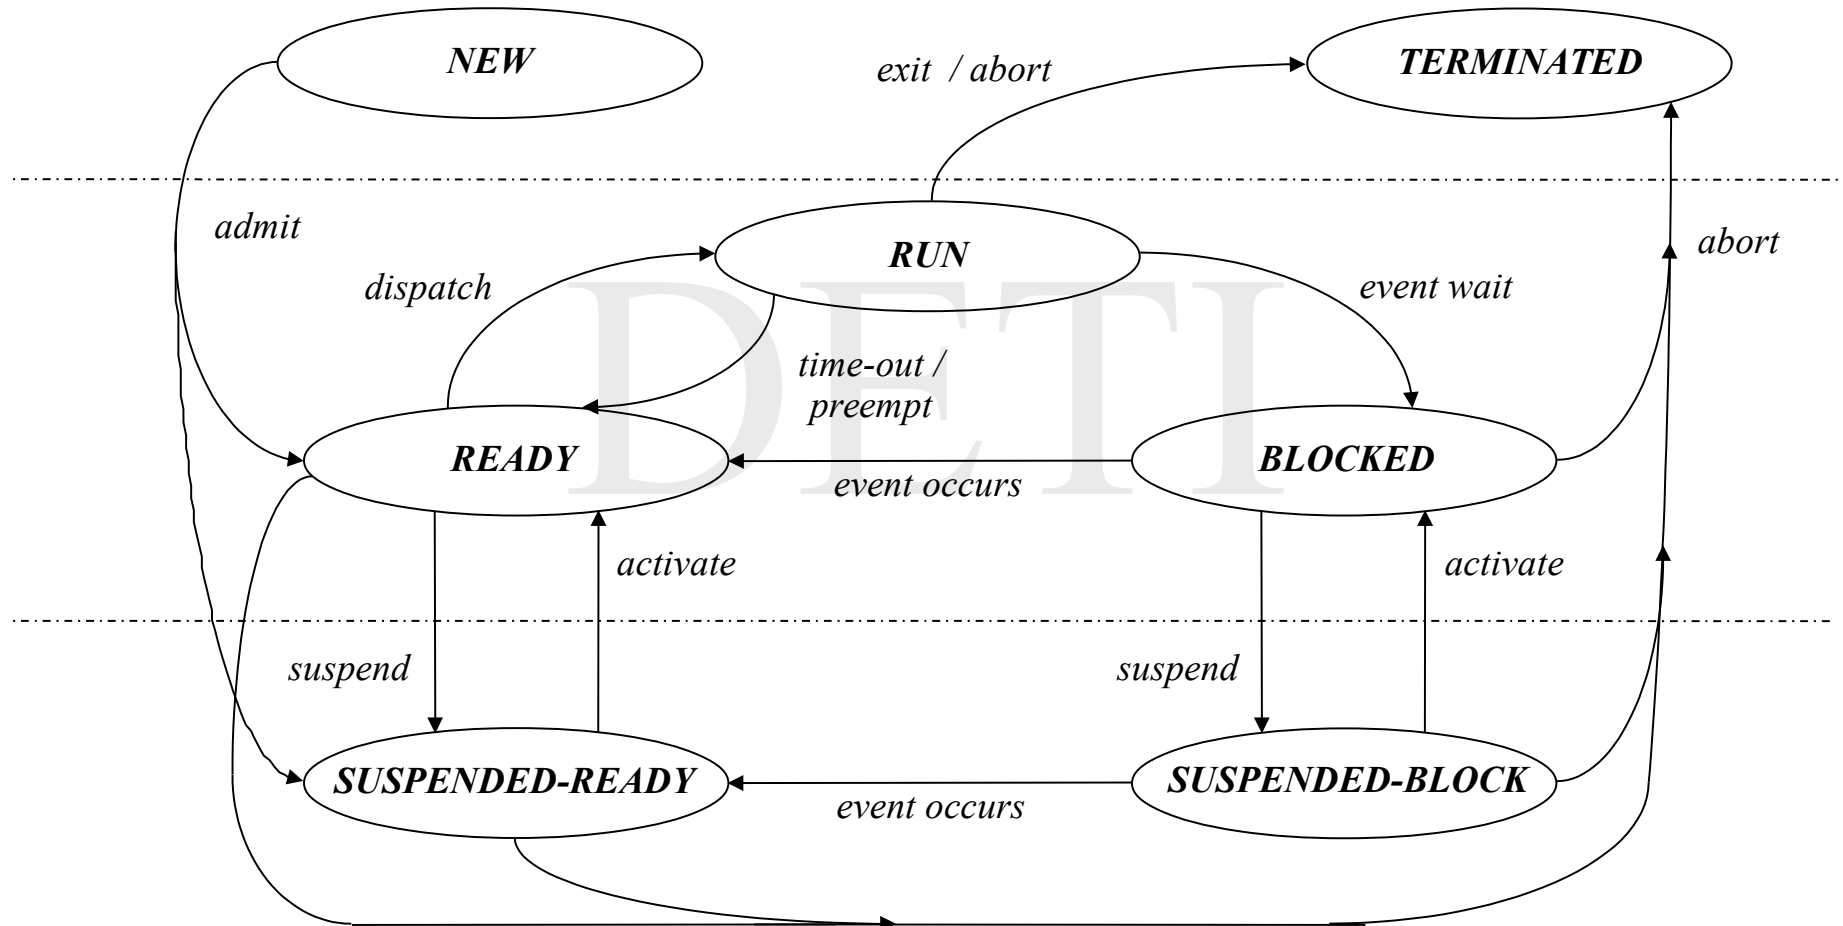
    status = EXIT_SUCCESS;
    pthread_exit (&status);
}

...
```

- there is not glibc wrapper for `gettid`
  - it has to be called indirectly, using system call `syscall`
- The TID of the main thread is the same as the PID of the process
  - actually, they are the same entity



# Process switching



## *Process switching*

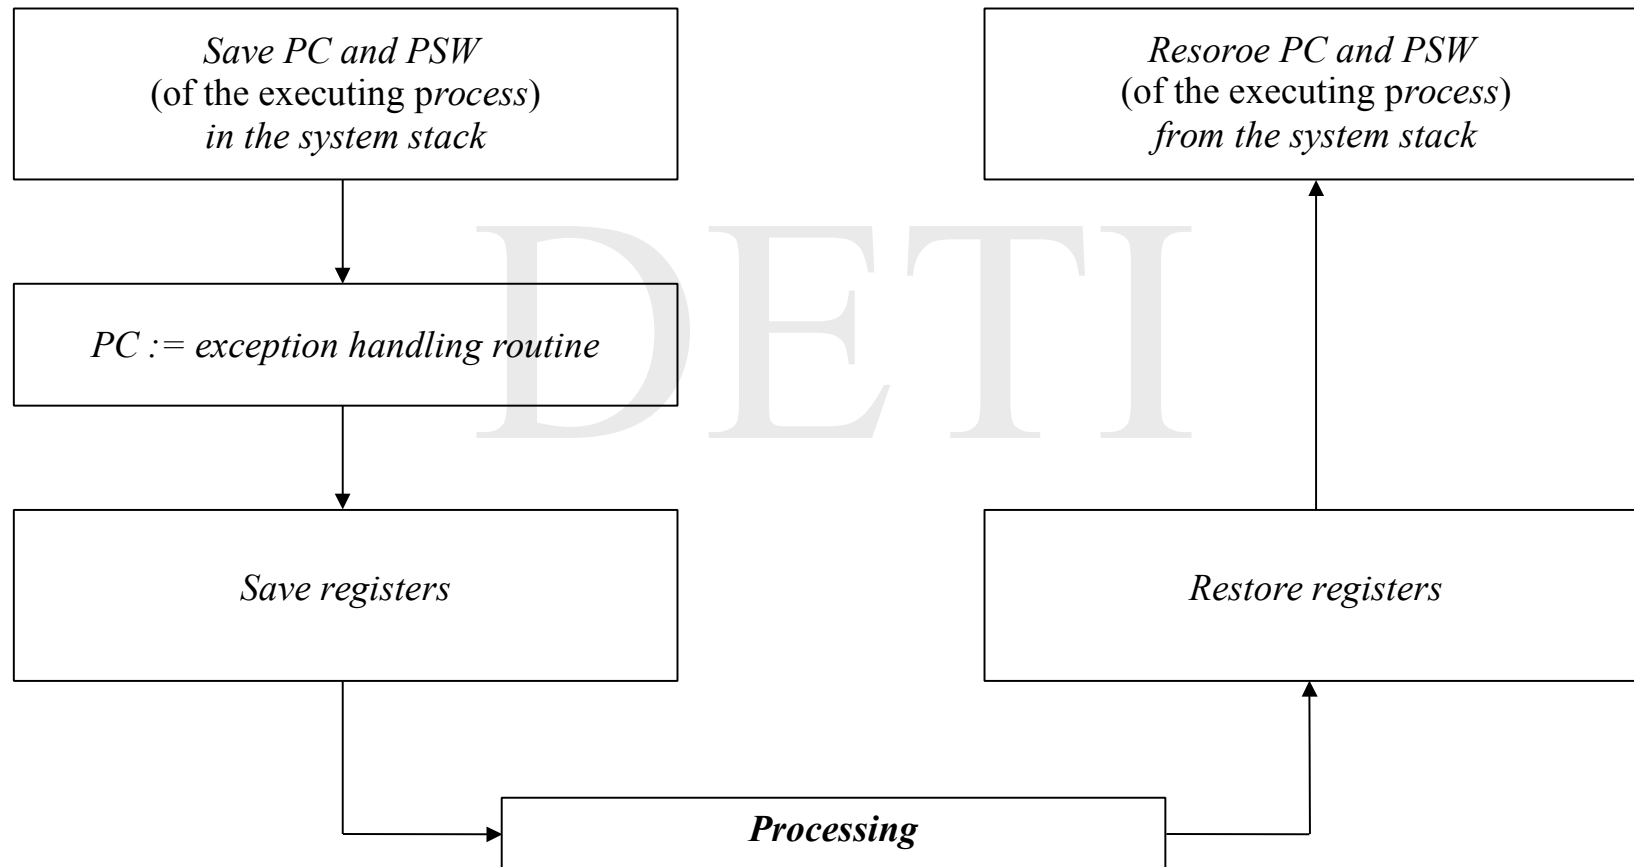
- ♦ Current processors have two functioning modes:
  - ♦ *supervisor mode* – all instruction set can be executed
    - ♦ is a privileged mode, reserved to the operating system
  - ♦ *user mode* – only part of the instruction set can be executed
    - ♦ input/output instructions are excluded as well as those that modify control registers
    - ♦ is the normal mode of operation
- ♦ Switching from user mode to supervisor mode is only possible through an *exception* (for security reasons)
- ♦ An exception can be caused by:
  - ♦ I/O interrupt
  - ♦ illegal instruction (division by zero, bus error)
  - ♦ trap instruction (software interruption)

## *Process switching*

- ♦ The operating system should function in supervisor mode
  - ♦ in order to have access to all the functionalities of the processor
- ♦ Thus kernel functions (including system calls) must be fired by
  - ♦ hardware (interrupt)
  - ♦ trap (software interruption)
- ♦ There is a uniform operation environment: *exception handling*
- ♦ Process switching can be seen the same way, with a small difference

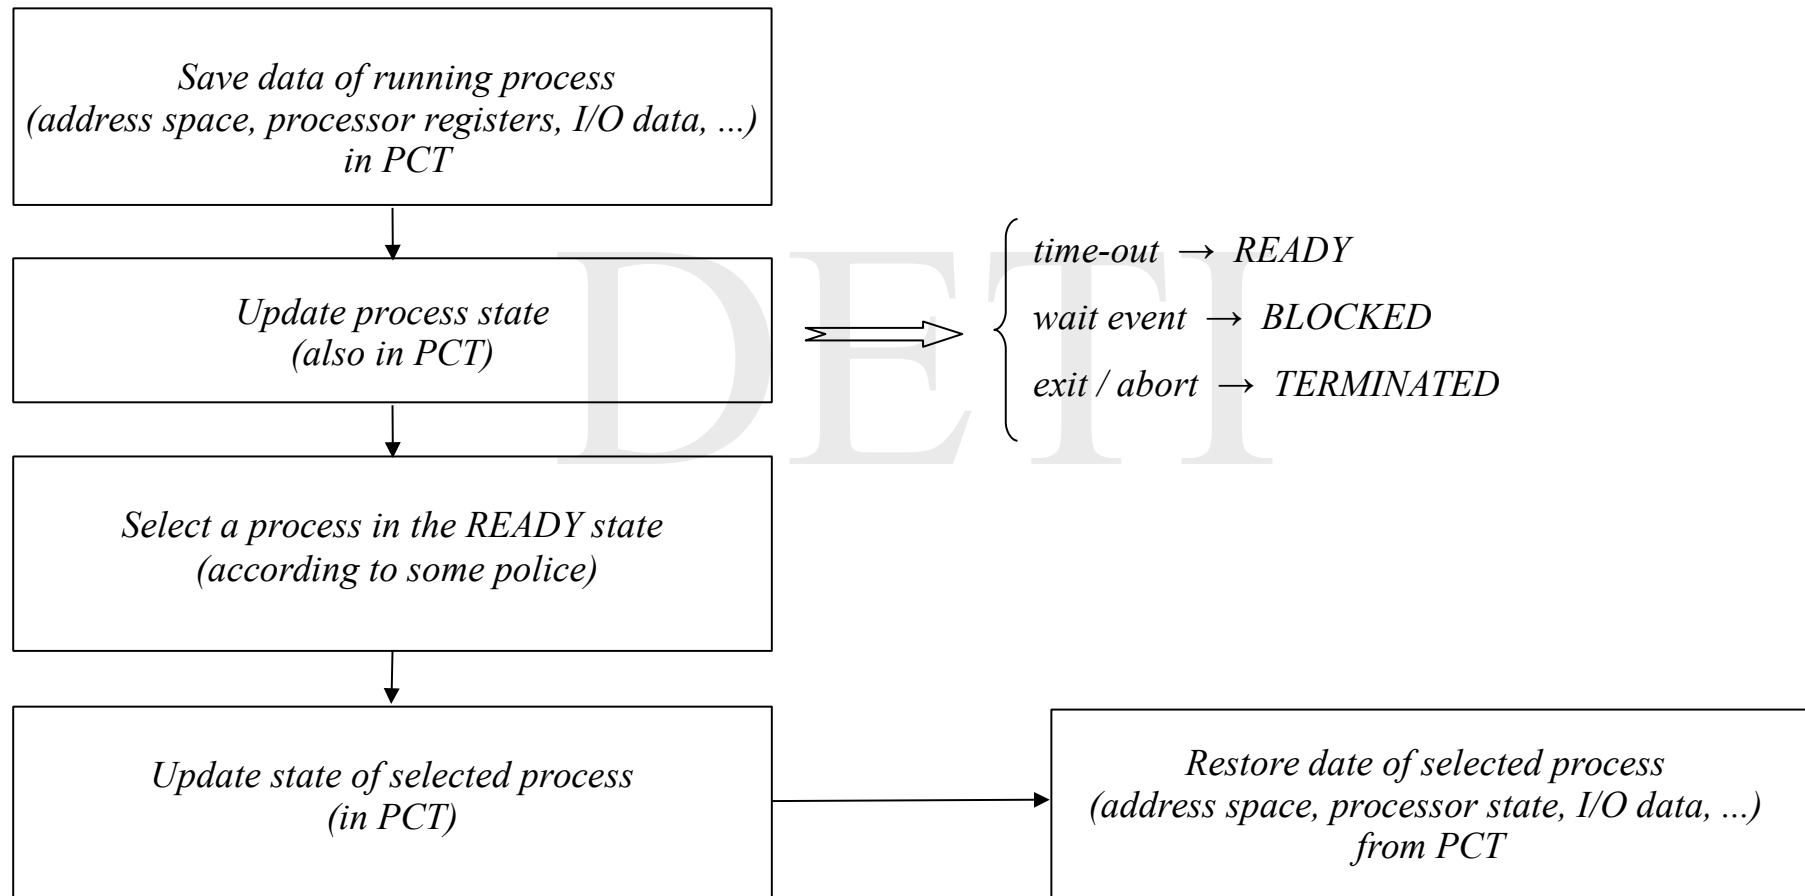
# *Process switching*

## *Processing a (normal) exception*



# Process switching

## Processing a process switching



## *Bibliography*

*Operating Systems Concepts*; Silberschatz, Galvin & Gagne; John Wiley & Sons, 9<sup>th</sup> Ed

- Chapter 3: *Processes* (sections 3.1 to 3.3)
- Chapter 4: *Threads* (sections 4.1 and 4.4.1)

*Modern Operating Systems*; Tanenbaum & Bos, Prentice-Hall International Editions, 4<sup>rd</sup> Ed

- Chapter 2: *Processes and Threads* (sections 2.1 and 2.2)

*Operating Systems*, W. Stallings, Prentice-Hall International Editions, 7<sup>th</sup> Ed

- Chapter 3: *Process Description and Control* (sections 3.1 to 3.5 and 3.7)
- Chapter 4: *Threads* (sections 4.1, 4.2 and 4.6)